

Alexander Karpekov

alex.karpekov@gmail.com | +1 (650) 660-7729 | U.S. (Permanent Resident) | LinkedIn: </karpekov> | alexkarpekov.com

Future: Work in AI and ML research & development space with a focus on model interpretability and explainability, and causal impact. Join a dynamic, synergistic team that values diverse strengths and has a shared vision for impact and success. Looking for in-office or hybrid opportunities, I believe in the importance of interpersonal connections. Preferred location: NYC.

Past: Data Scientist and ML researcher with 10 years of Big Tech and startup experience. Completed second Master's in CS while working full time at Google and being promoted twice. Studied and worked in 9 cities across 5 countries: San Francisco, New York, San Diego, Atlanta (U.S), Dublin (Ireland), London (U.K.), Strasbourg (France), Moscow, and St. Petersburg (Russia).

EDUCATION

Georgia Institute of Technology

2024 | Atlanta, Georgia | 3.9/4.0

MS in Computer Science

Machine Learning Specialization

Completed 2nd Master's Degree remotely while working full time at Google. Took a semester-long sabbatical to take classes and conduct research on Atlanta campus.

University of California, San Diego

2015 | San Diego, California | 3.8/4.0

MA in International Economics

Worked as a Teaching Assistant for 3 graduate-level classes in Statistics and Econometrics, leading sessions for 120+ students. Received the best TA award.

Moscow State University

2013 | Moscow, Russia | 92/100

BA in Political Science

SELECTED COURSEWORK

Deep Learning for Text
Generative Artificial Intelligence
Explainable Artificial Intelligence
Reinforcement Learning
Computer Vision
Advanced Econometrics
Survey Design and Evaluation
GIS and Spatial Data Analysis

SKILLS

Deep Learning: pytorch, tensorflow, keras, huggingface (transformers & datasets), langchain, lancedb, sklearn, statsmodels, xgboost

Python: numpy, pandas, scipy, jupyter, colab, matplotlib, altair, plotnine

Programming (in order of proficiency):

Python, SQL, R, Stata, C++, Java, Javascript

Languages: EN, RU, FR, DE, ZH

HOBBIES

Rowing (GeorgiaTech Crew), CrossFit (8 years), Snowboarding, Rock Climbing.

EXPERIENCE

GOOGLE | Senior Data Scientist @ YouTube and Google Search

2017 – Present | San Francisco, California and Dublin, Ireland

EXPERTISE: Statistical analysis using A/B testing and causal impact methodologies like propensity score matching; Clustering and classifications tasks; Embedding space construction; User data analytics and visualization, managing datasets with billions of entries.

Project Highlights

- Developed a pathfinding algorithm in song embedding space, improving music recommendations that led to 3% boost in user engagement and music discovery rates.
- Created a novel Transformer-based methodology to embed user action sequences on YouTube music app, which informed YouTube-wide product strategy for 2024.
- Implemented a new methodology to cluster YouTube multi-billion music corpus using text, sound, search, and co-watch embeddings, which led to a 30% reduction in harmful watchtime and a 0.5% increase in music revenue (\$100s millions).
- Created a new counterfactual causal impact methodology to evaluate the impact of the new feature launch on user engagement and conversion that helped establish no statistically significant long-term effects on key business metrics. The analysis was instrumental to halt the global rollout at Engineering and Product VP-level.

DATAMINR | Data Analyst

2015 – 2017 | New York, NY and London, U.K.

- Built statistical models to automatically classify Twitter user handles.
- Conducted Twitter user clustering and unsupervised learning using networks analysis methodologies to improve news discovery algorithms.
- Led company-wide effort for reporting automation using Python instead of Excel.

PROJECTS

Transformer Block Visualizer

2024 @ GT

[Ongoing] Developing a Transformer block interpretability tool for educators and ML developers that sheds light on the Transformer training process.

LLM Sensitivity to Prompt Perturbations

2023 @ GT

Created a statistical tool to evaluate LLM (GPT3, GPT4 and Llama 2) sensitivity to various prompt perturbations, such as lexical, semantic, structural, and tonal changes in the LLM input.

Using Stock Market Data to value Political Connections | 2015 @ UCSD

Created a methodology to estimate the impact of international sanctions on the Russian companies' market value using stock data and Capital Asset Pricing Model.

Gradient-Attention for Feature Importance Extraction in Transformers | 2023 @ GT

Created a feature extraction tool for Transformer-like models (e.g., BERT and BART) for sentiment analysis of text data.

Reinforcement Learning Model Interpretability | 2023 @ GT

Created a Q-value evaluation technique for Lunar lander RL environment to assess the agent learning progress throughout training.

Policy Impact Evals using Satellite Data 2014 @ UCSD

Used nightlight brightness data and GIS tools to approximate GDP growth in China. Ran geographically-weighted regression to evaluate a 15% increase in GDP as a result of the province-level policy in 2008-2013.